



3D Rendering and Ray Casting

Michael Kazhdan
(600.357 / 600.457)


HB Ch. 13.7, 14.6

FvDFH 15.5, 15.10



Rendering


- Generate an image from geometric primitives



Geometric
Primitives



Rendering




Raster
Image



Rendering


- Generate an image from geometric primitives



3D




Rendering



2D

3D Rendering Example




What issues must be addressed by a 3D rendering system?



Overview

- 3D scene representation
- 3D viewer representation
- What do we see?
- How does it look?






Overview

- 3D scene representation
- 3D viewer representation
- What do we see?
- How does it look?


How is the 3D scene described in a computer?






3D Scene Representation

- Scene is usually approximated by 3D primitives
 - Point
 - Line segment
 - Polygon
 - Polyhedron
 - Curved surface
 - Solid object
 - etc.



3D Point

- Specifies a location






3D Point

- Specifies a location
 - Represented by three coordinates
 - Infinitely small

```
class Point3D
{
public:
    Coordinate x;
    Coordinate y;
    Coordinate z;
};
```

• (x,y,z)





3D Vector

- Specifies a direction and a magnitude






3D Vector

- Specifies a direction and a magnitude
 - Represented by three coordinates
 - Magnitude $||V|| = \sqrt{dx^2 + dy^2 + dz^2}$
 - Has no location

```
class Vector3D
{
public:
    Coordinate dx;
    Coordinate dy;
    Coordinate dz;
};
```






3D Vector

- Specifies a direction and a magnitude
 - Represented by three coordinates
 - Magnitude $||V|| = \sqrt{dx\ dx + dy\ dy + dz\ dz}$
 - Has no location

```
class Vector3D
{
public:
    Coordinate dx;
    Coordinate dy;
    Coordinate dz;
};
```




- Dot product of two 3D vectors
 - $V_1 \cdot V_2 = dx_1 dx_2 + dy_1 dy_2 + dz_1 dz_2$
 - $V_1 \cdot V_2 = ||V_1|| ||V_2|| \cos(\Theta)$



3D Vector

- Specifies a direction and a magnitude
 - Represented by three coordinates
 - Magnitude $\|V\| = \sqrt{dx^2 + dy^2 + dz^2}$
 - Has no location

```
class Vector3D
{
public:
    Coordinate dx;
    Coordinate dy;
    Coordinate dz;
};
```



- Cross product of two 3D vectors
 - $V_1 \times V_2 = \text{Vector normal to plane } V_1, V_2$
 - $\|V_1 \times V_2\| = \|V_1\| \|V_2\| \sin(\Theta)$





Cross Product: Review

- Let $U = V \times W$:
 - $U_x = V_y W_z - V_z W_y$
 - $U_y = V_z W_x - V_x W_z$
 - $U_z = V_x W_y - V_y W_x$
- $V \times W = - W \times V$ (remember “right-hand” rule)
- We can do similar derivations to show:
 - $V_1 \times V_2 = \|V_1\| \|V_2\| \sin(\Theta) n$, where n is unit vector normal to V_1 and V_2
 - $\|V_1 \times V_1\| = 0$
- <http://physics.syr.edu/courses/java-suite/crosspro.html>



3D Line Segment

- Linear path between two points






3D Line Segment

- Use a linear combination of two points
 - Parametric representation:
 - » $P = P_1 + t (P_2 - P_1), \quad (0 \leq t \leq 1)$

```
class Segment3D
{
public:
    Point3D P1;
    Point3D P2;
};
```






3D Ray

- Line segment with one endpoint at infinity
 - Parametric representation:
 - » $P = P_1 + t V, \quad (0 \leq t < \infty)$

```
class Ray3D
{
public:
    Point3D P1;
    Vector3D V;
};
```






3D Line

- Line segment with both endpoints at infinity
 - Parametric representation:
 - » $P = P_1 + t V, \quad (-\infty < t < \infty)$


```
class Line3D
{
public:
    Point3D P1;
    Vector3D V;
};
```





3D Plane

- A linear combination of three points





3D Plane

- A linear combination of three points


- Implicit representation:

» $P \cdot N - d = 0$, or

» $ax + by + cz - d = 0$

```
class Plane3D
{
public:
    Vector3D N;
    Distance d;
};
```

- N is the plane “normal”
 - » Unit-length vector
 - » Perpendicular to plane
- d is the signed distance of the plane from the origin.





3D Polygon


- Area “inside” a sequence of coplanar points

- Triangle
- Quadrilateral
- Convex
- Star-shaped
- Concave
- Self-intersecting

```
class Polygon3D
{
public:
    Point3D *points;
    int npoints;
};
```

Points are in counter-clockwise order



- Holes (use > 1 polygon struct)





3D Sphere

- All points at distance “r” from point “(c_x, c_y, c_z)”
 - **Implicit representation:**
 - » $(x - c_x)^2 + (y - c_y)^2 + (z - c_z)^2 = r^2$
 - **Parametric representation:**
 - » $x = r \cos(\phi) \cos(\Theta) + c_x$
 - » $y = r \cos(\phi) \sin(\Theta) + c_y$
 - » $z = r \sin(\phi) + c_z$



```
class Sphere3D
{
public:
    Point3D center;
    Distance radius;
};
```



Other 3D primitives

- Cone
- Cylinder
- Ellipsoid
- Box
- Etc.



3D Geometric Primitives


- More detail on 3D modeling later in course
 - Point
 - Line segment
 - Polygon
 - Polyhedron
 - Curved surface
 - Solid object
 - etc.



Overview

- 3D scene representation
- 3D viewer representation
- What do we see?
- How does it look?


How is the viewing device described in a computer?





Camera Models

- The most common model is pin-hole camera
 - All captured light rays arrive along paths toward focal point without lens distortion (everything is in focus)



Other models consider ...

Depth of field


Motion blur

Lens distortion



Camera Parameters

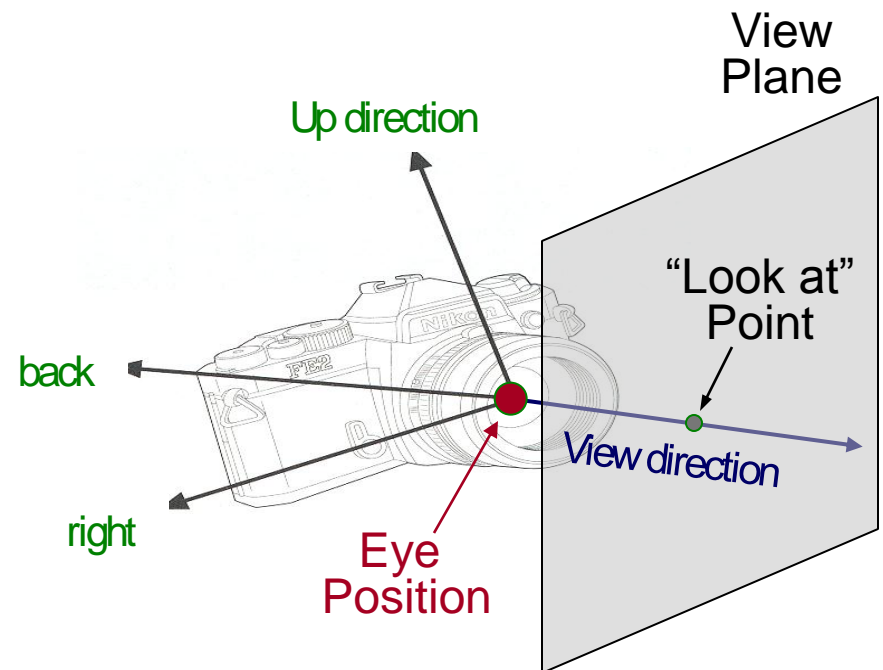
- What are the parameters of a camera?





Camera Parameters

- Position
 - Eye position (p_x, p_y, p_z)
- Orientation
 - View direction (d_x, d_y, d_z)
 - Up direction (u_x, u_y, u_z)
- Aperture
 - Field of view (x_{fov}, y_{fov})
- Film plane
 - “Look at” point
 - View plane normal



Other Models: Depth of Field



Close Focused



Distance Focused



Other Models: Motion Blur

- Mimics effect of open camera shutter
- Gives perceptual effect of high-speed motion
- Generally involves temporal super-sampling




Brostow & Essa



Traditional Pinhole Camera


- The film sits behind the pinhole of the camera.





Traditional Pinhole Camera


- The film sits behind the pinhole of the camera.
- Rays come in from the outside, pass through the pinhole, and hit the film plane.





Traditional Pinhole Camera


- The film sits behind the pinhole of the camera.
- Rays come in from the outside, pass through the pinhole, and hit the film plane.





Virtual Camera


- The film sits in front of the pinhole of the camera.





Virtual Camera


- The film sits in front of the pinhole of the camera.
- Rays come in from the outside, pass through the film plane, and hit the pinhole.





Virtual Camera


- The film sits in front of the pinhole of the camera.
- Rays come in from the outside, pass through the film plane, and hit the pinhole.





Overview

- 3D scene representation
- 3D viewer representation
- Ray Casting
 - What do we see?
 - How does it look?






Ray Casting

- Rendering model
- Intersections with geometric primitives
 - Sphere
 - Triangle
- Acceleration techniques
 - Bounding volume hierarchies
 - Spatial partitions
 - » Uniform grids
 - » Octrees
 - » BSP trees



Ray Casting


- We invert the process of image generation by sending rays out from the pinhole, and then we find the first intersection of the ray with the scene.





Ray Casting


- The color of each pixel on the view plane depends on the radiance emanating from visible surfaces





Ray Casting

- For each sample ...
 - Construct ray from eye position through view plane
 - Find first surface intersected by ray through pixel
 - Compute color sample based on surface radiance






Ray Casting

- Simple implementation:

```
Image RayCast(Camera camera, Scene scene, int width, int height)
{
    Image image = new Image(width, height);
    for (int i = 0; i < width; i++) {
        for (int j = 0; j < height; j++) {
            Ray ray = ConstructRayThroughPixel(camera, i, j);
            Intersection hit = FindIntersection(ray, scene);
            image[i][j] = GetColor(hit);
        }
    }
    return image;
}
```


- Where are we looking?
- What are we seeing?
- What does it look like?

Constructing a Ray Through a Pixel





Constructing a Ray Through a Pixel




The ray has to originate at P_0 , the position of the camera. So the equation for the ray is of the form:

$$\text{Ray} = P_0 + tV$$




Constructing a Ray Through a Pixel



If the ray passes through the point P ,
then the solution for V is:

$$V = (P - P_0) / \|P - P_0\|$$

Constructing a Ray Through a Pixel






Constructing Ray Through a Pixel

- 2D Example: Side view of camera at P_0
 - What is the position of the i -th pixel, $P[i]$?

θ = frustum half-angle (given), or field of view

d = distance to view plane (arbitrary = you pick)





Constructing Ray Through a Pixel


- 2D Example: Side view of camera at P_0
 - What is the position of the i -th pixel, $P[i]$?

θ = frustum half-angle (given), or field of view

d = distance to view plane (arbitrary = you pick)

$$P_1 = P_0 + d * \text{towards} - d * \tan(\theta) * \text{up}$$

$$P_2 = P_0 + d * \text{towards} + d * \tan(\theta) * \text{up}$$





Constructing Ray Through a Pixel

- 2D Example: Side view of camera at P_0
 - What is the position of the i -th pixel, $P[i]$?


θ = frustum half-angle (given), or field of view

d = distance to view plane (arbitrary = you pick)

$$P_1 = P_0 + d \cdot \text{towards} - d \cdot \tan(\theta) \cdot \text{up}$$

$$P_2 = P_0 + d \cdot \text{towards} + d \cdot \tan(\theta) \cdot \text{up}$$

$$\begin{aligned} P[i] &= P_1 + ((i+0.5)/\text{height}) \cdot (P_2 - P_1) \\ &= P_1 + ((i+0.5)/\text{height}) \cdot 2 \cdot d \cdot \tan(\theta) \cdot \text{up} \end{aligned}$$





Constructing Ray Through a Pixel

- 2D Example:
 - The ray passing through the i -th pixel is defined by:

$$\text{Ray} = P_0 + tV$$

- Where:


- P_0 is the camera position
- V is the direction to the i -th pixel:
$$V = (P[i] - P_0) / \|P[i] - P_0\|$$
- $P[i]$ is the i -th pixel location:

$$P[i] = P_1 + ((i+0.5)/\text{height}) * (P_2 - P_1)$$

- P_1 and P_2 are the endpoints of the view plane:

$$P_1 = P_0 + d * \text{towards} - d * \tan(\theta) * \text{up}$$


$$P_2 = P_0 + d * \text{towards} + d * \tan(\theta) * \text{up}$$



Constructing Ray Through a Pixel



- Figuring out how to do this in 3D is assignment 2





Ray Casting

- Simple implementation:

```
Image RayCast(Camera camera, Scene scene, int width, int height)
{
    Image image = new Image(width, height);
    for (int i = 0; i < width; i++) {
        for (int j = 0; j < height; j++) {
            Ray ray = ConstructRayThroughPixel(camera, i, j);
            Intersection hit = FindIntersection(ray, scene);
            image[i][j] = GetColor(hit);
        }
    }
    return image;
}
```



Ray Casting

- Simple implementation:

```
Image RayCast(Camera camera, Scene scene, int width, int height)
{
    Image image = new Image(width, height);
    for (int i = 0; i < width; i++) {
        for (int j = 0; j < height; j++) {
            Ray ray = ConstructRayThroughPixel(camera, i, j);
            Intersection hit = FindIntersection(ray, scene);
            image[i][j] = GetColor(hit);
        }
    }
    return image;
}
```



Ray-Scene Intersection


- Intersections with geometric primitives
 - Sphere
 - Triangle
- Acceleration techniques
 - Bounding volume hierarchies
 - Spatial partitions
 - » Uniform (Voxel) grids
 - » Octrees
 - » BSP trees



Ray-Sphere Intersection

Ray: $P = P_0 + tV$

Sphere: $|P - O|^2 - r^2 = 0$






Ray-Sphere Intersection

Ray: $P = P_0 + tV$

Sphere: $|P - O|^2 - r^2 = 0$

Substituting for P , we get:

$$|P_0 + tV - O|^2 - r^2 = 0$$





Ray-Sphere Intersection

Ray: $P = P_0 + tV$

Sphere: $|P - O|^2 - r^2 = 0$

Substituting for P , we get:

$$|P_0 + tV - O|^2 - r^2 = 0$$

Solve quadratic equation:


$$at^2 + bt + c = 0$$

where:

$$a = 1$$

$$b = 2 V \cdot (P_0 - O)$$

$$c = |P_0 - O|^2 - r^2 = 0$$





Ray-Sphere Intersection

Ray: $P = P_0 + tV$

Sphere: $|P - O|^2 - r^2 = 0$

Substituting for P , we get:

$$|P_0 + tV - O|^2 - r^2 = 0$$

Solve quadratic equation:


$$at^2 + bt + c = 0$$

where:

a
b
c

Generally, there are two solutions to the quadratic equation, giving rise to points P and P' .

You want to return the first hit.






Ray-Sphere Intersection

- Need normal vector at intersection for lighting calculations

$$N = (P - O) / ||P - O||$$






Ray-Scene Intersection

- Intersections with geometric primitives
 - Sphere
 - » Triangle
- Acceleration techniques
 - Bounding volume hierarchies
 - Spatial partitions
 - » Uniform grids
 - » Octrees
 - » BSP trees



Ray-Triangle Intersection

- First, intersect ray with plane
- Then, check if point is inside triangle





Ray-Plane Intersection

Ray: $P = P_0 + tV$

Plane: $P \cdot N - d = 0$


Algebraic Method

Substituting for P , we get:

$$(P_0 + tV) \cdot N - d = 0$$

Solution:

$$t = -(P_0 \cdot N - d) / (V \cdot N)$$





Ray-Triangle Intersection I

- Check if point is inside triangle algebraically

For each side of triangle

$$V_1 = T_1 - P_0$$


$$V_2 = T_2 - P_0$$

$$N_1 = V_2 \times V_1$$

$$\text{if } ((P - P_0) \cdot N_1 < 0)$$

return FALSE;

end





Ray-Triangle Intersection II

- Check if point is inside triangle parametrically


Every point P inside the triangle can be expressed as the weighted average of the corners:

$$P = \alpha T_1 + \beta T_2 + \gamma T_3$$

where:

$$0 \leq \alpha, \beta, \gamma \leq 1$$

$$\alpha + \beta + \gamma = 1$$





Ray-Triangle Intersection II

- Check if point is inside triangle parametrically

Solve for α, β, γ such that:

$$P = \alpha T_1 + \beta T_2 + \gamma T_3$$

And

$$\alpha + \beta + \gamma = 1$$

Check if point inside triangle.

$$0 \leq \alpha, \beta, \gamma \leq 1$$

